Eight philosophy candidates

Originality audit and development drafts

22 October 2026 · Prepared by Codex

The earlier shortlist was too generous. These are eight developed arguments, not eight established discoveries. The closer reading supports strong demotions, one merger, and several conditional prospects.

1. Revocation: who pays for the independent checking that makes delegated authority controllable?

2. Models and accountability: what must a model establish before an institution uses it to blame someone?

3. Preserving correction: may an institution shift the future cost of checking its assurances onto those who relied on them?

4. Inherited authority: when must present trustees justify continuing to bind successors to a founder's purposes?

5. Translating disagreement: who must correct an institutional account that misclassifies what someone is disputing?

6. Forgotten obligations: can a person pressing a surviving claim owe help making its basis accessible?

7. Reopening inquiry: what is owed when the evidence needed to reopen a case is accessible only after reopening?

8. Contesting objectives: who is responsible when specialized complaint procedures leave a justified challenge with no forum?

Development order: C045 first; C004 and C085 together second; C063 and C052 conditionally thereafter. Do not submit C013 or C021's original broad claims as discoveries. C111 needs a narrower contribution.

The eight drafts total approximately 7,850 words before references. They are development essays, not completed 6,000-word contest submissions. The final pages identify inaccessible texts. A separate audit and complete access ledger accompany this PDF in the GitHub package.

A revocation right whose evidence belongs to the delegate

Candidate C013 · Development draft · 22 October 2026

Originality verdict: The broad independence thesis is already substantially occupied. The argument below is a narrower proposal about the allocation of verification costs. It is not a certified new result.

The problem

Imagine that a cooperative delegates procurement to an exceptionally capable manager. Members can dismiss her immediately. They also receive accurate weekly reports. Nevertheless, every report is assembled from categories she defines, and only she can authorize an investigation into what the categories exclude. Nothing currently indicates misconduct. Is the cooperative's authority over its manager secure?

An easy answer says that dismissal remains legally available. A more demanding answer asks whether members have a usable way of discovering that dismissal is warranted. But this second answer is not our discovery. Cohen and Suzor's Contesting the public interest in AI governance (2024), especially §3.3, already connects meaningful contestation with independent information and examines information monopolies and external scrutiny. Relabeling that requirement as “epistemic revocability” would not establish originality.

The remaining question is distributive. Suppose independent investigation exists, but exercising it requires every member to become a procurement expert. Who must bear the cost of maintaining a practical capacity to reconsider delegation?

A narrower proposal

When an institution represents a delegation as subject to continuing control, and the delegation foreseeably concentrates access to information needed to exercise that control, the arrangement should normally finance an adequate means of independent checking. Providing an abstract permission to purchase an audit does not discharge this responsibility when the audit's cost predictably defeats the promised control.

This is a claim about the responsibilities created by an institutional arrangement. It is not a claim that every delegate must personally pay for unlimited scrutiny. A small voluntary club might share the costs explicitly. A powerful organization governing dependent members might have to absorb them. The relevant question is whether the people asked to authorize the arrangement are offered practically usable control on terms they can fairly be expected to meet.

Compare two versions of the cooperative. In the first, the common budget pays for a periodically selected independent auditor, and members can trigger a limited supplementary inquiry at a manageable threshold. In the second, each dissenter must privately purchase an equally good audit. Assume the same manager, evidence, performance, and formal dismissal procedure. An affluent member could investigate in either version. Ordinary members cannot investigate in the second.

The difference is not a difference in the manager's trustworthiness. Nor does it establish that she lacks legitimate authority in every transaction. It concerns the distribution of the capacity that supposedly makes her continuing authority answerable to members. The second arrangement turns collective control into an option exercised mainly by the wealthy.

Why the duty follows

The proposed argument has three steps. First, a right offered as a safeguard cannot do its justificatory work if its foreseeable exercise conditions exclude its intended holders. Second, delegation can itself generate these conditions by concentrating expertise, records, and practical familiarity in the delegate. Third, the institution that deliberately creates that dependence has a stronger responsibility to arrange proportionate checking than an unrelated outsider has to subsidize any investigation whatsoever.

The controversial step is the third. It might be defended as a principle against imposing the costs of an institutional safeguard on those whose dependence makes the safeguard necessary. The argument therefore does not infer a positive duty merely from the desirability of an audit. It

connects the duty to an organization's own claim that its concentration of authority remains controllable.

This also explains why independence is a matter of powers rather than labels. An auditor paid from the common budget can be independent if the manager cannot suppress findings or dismiss the auditor at will. Conversely, an outside contractor can be captured if the manager determines every question the contractor may ask. Separate employment is neither necessary nor sufficient for the proposed protection.

Limits and objections

The strongest objection is that fiduciary duties, republican freedom, and fair access to institutional remedies already entail all of this. That objection is serious. The unavailable DeMott and Pettit texts listed in the access register are necessary comparisons. Until they are read, this draft should be treated as a possible synthesis or application, not a new principle.

A second objection is that checking defeats the point of delegation. Members delegate because independent verification is expensive. The answer cannot be that every decision requires duplication. Adequacy should depend on stakes, reversibility, existing public information, and the reliability of alternative checks. Sampling, pooled expertise, and protected whistleblowing might be enough. Some low-stakes arrangements need no standing audit at all.

A third objection concerns confidentiality. Independent checking need not mean public release of every record. A trusted intermediary can verify a disputed question without exposing irrelevant personal information. The required power is to investigate and communicate a warranted challenge through an effective route.

What would make this worth an essay?

Its philosophical significance would lie in showing that maintaining revocable authority creates a distinctive responsibility to distribute the costs of knowing when to revoke. But the case is not yet made that this responsibility differs from familiar fairness and fiduciary principles. The appropriate next test is a case in which those principles appear satisfied while verification costs still undermine the claimed authority. If no such case survives, retain this as a useful institutional application and remove it from the originality shortlist.

References and access

- Cohen, Tegan, and Nicolas P. Suzor. 2024. “Contesting the public interest in AI governance.” Internet Policy Review 13(3). Source. Full text available; relevant information-independence discussion read.
- DeMott, Deborah A. 2014. “The Fiduciary Character of Agency and the Interpretation of Instructions.” Source. Repository description only; full text requested.
- Pettit, Philip. 2000. “Democracy, Electoral and Contestatory.” Designing Democratic Institutions, pp. 105-144. Source. Publication verified; linked full text inaccessible.

What an accountability model must preserve

Candidate C052 · Development draft · 22 October 2026

Originality verdict: Preserving strategic ability under abstraction is established technical work. A narrower philosophical argument about the evidential burden of using a model to blame someone remains worth testing. The example below is an illustration, not a new theorem.

The change of question

A model can answer “What intervention would have prevented this outcome?” without answering “Could this person reasonably have performed that intervention?” This distinction matters whenever an organization converts a useful explanation of failure into a judgment against a particular person.

It is tempting to announce that causal abstraction therefore needs a new requirement to preserve responsibility. That announcement would be premature. Beckers and Halpern's Abstracting Causal Models (2019) already distinguishes demanding forms of abstraction by what they preserve. A separate technical literature explicitly models agents' actions and information. Belardinelli and Lomuscio (2017), for example, define uniform strategies over an agent's local states and develop abstractions for strategic claims under imperfect information. A requirement to preserve agent-specific abilities is not an untouched idea.

The philosophical question is instead one of permitted inference: what must an institution establish before treating the success of a causal model as evidence of an individual's culpable failure?

A small example

A technician must choose either the left or the right switch. Exactly one prevents a shutdown. A hidden configuration determines which one. The technician's display looks identical in both configurations, and there is no available test before the deadline. There are no relevant earlier opportunities to improve the display, no negligence in becoming ignorant, and no independent duty to choose one side rather than the other.

In configuration L, choosing left prevents the shutdown. In configuration R, choosing right does. A retrospective investigation identifies the actual configuration and correctly concludes that the other action would have prevented the actual failure.

The investigation has identified a successful counterfactual intervention. It has not shown that the technician had an information-sensitive way of selecting it. Every configuration admits a successful action; no single action succeeds across the configurations the technician cannot distinguish. Those statements differ in the order of “for every configuration” and “there is an action.” That elementary distinction is familiar, and the uniform-strategy requirement in the cited technical work formalizes the relevant constraint.

Now compare a second workplace. Its physical switch mechanism and intervention responses are identical, but the technician has a reliable display of the configuration. The intervention table alone is unchanged; the person's capacity to choose appropriately is different. A model omitting the display might remain excellent for redesigning the mechanism. Its adequacy for that purpose does not establish its adequacy for a finding of knowing or negligent noncompliance.

Neither example proves that alternative possibilities are necessary for every kind of moral responsibility. Nor does lack of a guaranteed strategy automatically excuse risky conduct. The examples isolate one particular accusation: that the person culpably failed to choose the known effective intervention. That accusation requires facts the intervention table does not supply.

A proposed rule for institutional use

An institution invoking a model to support a specified accusation should identify the normative premises of that accusation and show that the evidence preserves or independently establishes the facts those premises require. A model that establishes physical preventability cannot silently acquire the missing premise that this agent could recognize and execute the preventive action.

This requirement is deliberately local. A model need not contain a complete theory of responsibility. If the accusation concerns a technician's failure to follow an intelligible warning, the institution needs evidence about access to the warning and capacity to act on it. If the accusation concerns negligent maintenance of the warning system, the relevant agent and time shift. If it concerns strict liability rather than fault, the argument must say so instead of borrowing the rhetoric of culpability.

The duty also concerns who bears uncertainty. Suppose the organization designed the display, kept the configuration records, and later destroyed the records of what the technician could see. It should not infer adequate awareness simply because its remaining model explains the physical mechanism. Whether destruction warrants a stronger presumption against the organization requires an additional argument about evidence and fairness; this draft does not smuggle that conclusion into the definition of abstraction.

Why this is more than asking for a bigger model

The proposal does not demand maximal detail. It asks the institution to state the claim for which the model is evidence. An abstraction that omits personal information might be preferable for planning repairs and unsuitable for assigning individual fault. The same omission can thus be harmless or decisive depending on the conclusion drawn.

This is also why successful validation on predictions is not a universal certificate. Predictive accuracy supports predictions in the tested domain. It does not by itself validate an inference that depends on an omitted informational or normative relation. There may be no predictive error to expose: the error can occur when the institution changes its question without changing its evidential warrant.

The originality test

The technical ingredients are inherited. The remaining possibility is an account of institutional responsibility for transferring a model between explanatory and accusatory purposes, including how the burden of establishing adequacy should be allocated. That could matter for automated workplace discipline and accident investigation, but a new application is not enough for a contest contribution.

A serious next draft must compare this proposal with existing work on model adequacy, epistemic conditions of responsibility, and evidential burdens. It must find a disputed case on which the proposed allocation changes a defensible verdict. If it merely restates “use a model appropriate to the question,” the idea should be rejected as too general. This audit therefore demotes its earlier priority rather than dressing the familiar technical distinction as an original formal result.

References and access

- Beckers, Sander, and Joseph Y. Halpern. 2019. “Abstracting Causal Models.” AAAI. Source. Relevant definitions and examples read; not every proof independently checked.
- Belardinelli, Francesco, and Alessio Lomuscio. 2017. “Agent-based Abstractions for Verifying Alternating-time Temporal Logic with Imperfect Information.” AAMAS, pp. 1259-1267. Source. Definitions 1-4 and account of the abstraction result read; proof not independently verified.
- Jamroga, Wojciech, and Yan Kim. 2023. “Practical Abstraction for Model Checking of Multi-Agent Systems.” KR, pp. 384-394. Source. Introduction and related-work map used to locate the closer 2017 source.

The cost of making a claim uncheckable

Candidate C004 · Development draft · 22 October 2026

Originality verdict: Generic duties to preserve evidence or remain corrigible are not new. The narrower candidate concerns transferring the cost of future correction when an institution withdraws a checking route on which it has induced reliance. Develop together with C085.

A justified conclusion and a later decision

A public laboratory certifies that a construction material passes a specified safety test. Its evidence is currently good. Builders reasonably rely on the certification, and retaining every physical sample forever would be expensive. Five years later the laboratory proposes to destroy the remaining samples. No new reason to doubt the certification has appeared.

One question is whether the laboratory should believe its original conclusion. Another is whether it should conduct a new test now. A third is whether it may eliminate the only affordable means by which other people could investigate an important future challenge. The answers need not coincide.

The proposed thesis concerns the third question. An institution can be justified in retaining its belief and in ending active inquiry while owing preservation, replacement, or fair transfer of a route for correction. The duty arises here from continuing, foreseeable reliance together with control over an otherwise inaccessible means of checking. It does not follow merely from the abstract possibility of error.

Two ways of closing the archive

Suppose the laboratory can retain a representative set of samples at modest cost. Alternatively, it can destroy them and leave every building owner to purchase invasive testing of installed material if doubt later emerges. That second route is physically possible but prohibitively expensive for many owners. The proposed destruction therefore does more than save the laboratory money: it transfers the prospective cost of correcting its own authoritative assurance onto those who relied on it.

Now change the case. The laboratory can send equivalent samples and adequate records to an independent archive with guaranteed access. Destruction of its own copies no longer creates the same problem. What matters is the checking capacity, not the continued existence of every original object.

This comparison suggests a replacement principle. When an institution permissibly retires a means of checking a claim on which it has induced substantial continuing reliance, it should normally preserve an adequate substitute or fairly provide for the additional burdens its retirement creates. Whether it must do so depends on the stakes, cost, availability of alternatives, duration of reliance, and competing claims such as privacy.

Why reliance matters

There are many things one could learn more about. No person or institution can preserve every path to every possible discovery. Reliance narrows the duty because the institution has made its conclusion part of other people's practical arrangements. Those people may have forgone independent testing precisely because the institution occupied an authoritative role.

The principle is strongest when the assurance invited such dependence, the reliance was reasonable, the consequences remain significant, and the institution is uniquely positioned to maintain the route. It is weaker for casual conversation, trivial claims, or reliance the speaker neither invited nor could reasonably foresee. It may end when the relevant reliance ends. An institution cannot necessarily end it by announcing that it no longer wishes to be responsible while people remain trapped in the arrangements it encouraged.

The duty is best understood as moral and institutional, with an epistemic object. Calling it distinctively epistemic would require a further argument. Nothing here depends on proving that all responsibilities concerning inquiry arise from an autonomous epistemic norm.

Existing work and the actual gap

Flores and Woodard's Epistemic norms on evidence-gathering (2023) already challenges a restriction of epistemic evaluation to responses to evidence presently held. Longino's How Values Can Be Good for Science (2004), pp. 133-135, already connects responsible inquiry with effective criticism and recognizes later reassessment. Neither broad expansion is available as our contribution.

The closest unresolved obstacle is Porat and Stein's Liability for Uncertainty: Making Evidential Damage Actionable (1997). Its abstract and publication record concern responsibility for damaging another's evidential position. I could not obtain the full text. The proposed replacement principle may turn out to be a prospective application of its existing argument. That would change the novelty assessment substantially.

The possible increment is thus narrowly specified: a duty can arise before any dispute, without negligent destruction or a present duty to collect additional evidence, because retiring an institutionally supplied checking route reallocates the cost of future correction. A useful essay must show why existing evidential-damage and reliance principles do not already settle precisely that case.

Hard cases

One objection is that the probability of future error may be minute. The response is proportionality, not an unlimited preservation command. A cheap safeguard against a grave, durable risk may be warranted; an expensive safeguard against a negligible inconvenience may not be. No numerical threshold is supplied here, and the draft should not pretend to have solved that allocation problem.

A second objection concerns confidential records. An archive can itself harm people. Preserving correction may require a different route: controlled access, independent testing, or a restricted record sufficient for a particular challenge. Sometimes the privacy cost defeats preservation altogether. The principle then requires an honest account of the lost checking capacity rather than a fictional claim that nothing epistemically significant has changed.

A third objection is that induced reliance can be unreasonable. Someone who treats a clearly limited exploratory result as an unconditional guarantee does not automatically acquire a perpetual claim to laboratory resources. The assurance's content and reasonable use must be established independently.

What survives

The worthwhile proposal is not “keep evidence.” It is a responsibility for shifting the future cost of checking an authoritative claim. C085 asks what happens once a challenge is raised; this draft asks what capacity must remain available beforehand. Those are complementary stages of one argument. They should not be counted as two independent discoveries.

References and access

- Flores, Carolina, and Elise Woodard. 2023. “Epistemic norms on evidence-gathering.” Philosophical Studies 180:2547-2571. Source. Relevant opening arguments read.
- Longino, Helen E. 2004. “How Values Can Be Good for Science.” In Science, Values, and Objectivity, pp. 127-142. Source. Especially pp. 133-135 read.
- Porat, Ariel, and Alex Stein. 1997. “Liability for Uncertainty: Making Evidential Damage Actionable.” Cardozo Law Review 18:1891-1960. Source. Abstract/metadata only; full text requested.

Inheriting benefits without inheriting the founder's reasons

Candidate C111 · Development draft · 22 October 2026

Originality verdict: Intergenerational revisability is an established problem. A narrower argument about the justificatory burden of continuing a founder's mission is developed here, but its originality remains uncertain and a recent primary text is missing.

An inheritance can contain a demand

Imagine an educational foundation endowed for two centuries. Its founder selected a reasonable mission: provide technical training in communities neglected by existing institutions. A century later the foundation still provides substantial benefits. Its successors, however, have good reasons to prefer a different mixture of technical, artistic, and civic education. The trust's decision procedures instruct them to treat fidelity to the founder's original ranking as a decisive reason against revision.

The issue is not whether founders may ever attach conditions to gifts. Nor is it whether every inherited restriction is unjust. The question is whether the continuing benefit supplied by an inheritance is sufficient to justify treating its original evaluative standard as authoritative for people who did not choose it.

The initial candidate proposed to hold both welfare and the entire menu of feasible options fixed while varying revising authority. That comparison was too loose. A power to revise standards changes at least some counterfactual options. The defensible comparison holds current provision and independently specified material opportunities fixed, while allowing the institutional powers under examination to differ. It cannot establish that authority matters independently of every possible conception of opportunity.

The narrower thesis

An institution that continues to impose an inherited mission on dependent participants owes them a present justification for the mission's binding status. The fact that the mission originated in a beneficial gift can be relevant to that justification. It is not invariably sufficient. Where the institution also controls access to an important good, it may owe participants a usable procedure for proposing and assessing revisions that the founder would have rejected.

This is deliberately a claim against present institutional agents. Beckman's Democratic Legitimacy and Decisions for the Future (2025) challenges the inference from affecting future people to governing them and emphasizes the role of current authority. The proposed argument should accept that challenge. A founder's historical influence does not automatically make the founder a current ruler. Living trustees can nevertheless exercise authority badly when they treat a founder's preference as exempt from present justification.

Consider two foundations offering identical courses today. One requires trustees periodically to ask whether the mission remains justified to the people it serves. The other permits changes only when they can be interpreted as faithful to the founder's original purposes. Suppose a proposed change would protect the core educational benefit while responding to a serious unanticipated need. The second foundation cannot answer the proposal merely by repeating that the founder valued something else. Its trustees need a reason why that historical preference should settle the present question.

Why the historical source is not enough

One argument for fidelity appeals to donor expectations. Permitting revision could reduce future gifts. This is a real consideration, but it is an instrumental argument about sustaining a valuable practice. Its force varies with the likely consequences and the importance of the restriction. It does not establish an unlimited authority of the dead.

Another argument appeals to ownership. The founder could have consumed the resources, so successors cannot complain about receiving them on conditions. That argument is strongest for an optional gift among plentiful alternatives. It weakens when the institution becomes a durable provider of an essential good, when public cooperation sustains its position, or when later

participants cannot reasonably avoid its rules. An initially optional transaction can become part of a structure of dependence for people who never accepted its terms.

A third argument appeals to continuity. Projects need stable purposes, and periodic revision can destroy their identity. Stability can justify procedural thresholds and a presumption of continuity. It need not justify making inherited standards unrevisable in principle. A process can demand substantial reasons for change without demanding that every reason be expressible in the founder's own evaluative vocabulary.

What the literature already supplies

Hohnerlein's Democratic Revisability and the Legitimacy of Constitutional Entrenchments (2024) directly examines the tension between inherited constraints and later democratic revision, distinguishing interpretation from democratic power to change rules. The general insight that beneficial constraints can still raise authority problems is therefore not novel.

The remaining domain is narrower: institutions whose authority comes partly from an inherited private project but whose continuing role creates public or relational responsibilities. The candidate would need to explain when dependence changes the justificatory status of founder fidelity. This is likely to intersect with existing work on charitable trusts, donor intent, dead-hand control, and political legitimacy. Those literatures have not been exhaustively checked here.

Objections that matter

The danger of paternalism is immediate. A procedure that allows successors to rewrite every mission might simply let the powerful appropriate resources intended for a neglected group. Revision must therefore consider the interests that the original restriction protected. “Future people should decide” is not sufficient when some future people dominate others.

Nor does revisability require unanimity or a personal veto. It requires an appropriate opportunity to challenge, reasons responsive to the challenge, and a decision procedure with authority to revise. The correct constituency may include beneficiaries, workers, surrounding communities, and others affected by the institution's privileged position.

The strongest originality objection is that this is a familiar application of democratic revisability to endowments. That may be correct. A full comparison with André Santos Campos's 2026 Decoupling Generational Sovereignty is also still outstanding because the full text was inaccessible. The current draft is therefore a research prospectus with a developed argument, not a claim to have discovered a neglected dimension of intergenerational justice.

References and access

- Beckman, Ludvig. 2025. “Democratic Legitimacy and Decisions for the Future.” Res Publica. Source. Main argument read.
- Hohnerlein, Jakob. 2024. “Democratic Revisability and the Legitimacy of Constitutional Entrenchments: A Content-Specific Approach.” German Law Journal 25:726-749. Source. Introduction and revisability/interpretation discussion read.
- Santos Campos, André. 2026. “Decoupling Generational Sovereignty: The Cross-Temporal Role of the Sovereign People.” Moral Philosophy and Politics, advance online publication. Source. Abstract only; full text requested.

A translation can preserve the sentences and lose the dispute

Candidate C063 · Development draft · 22 October 2026

Originality verdict: The general importance of pragmatic meaning and power in translation is established. A limited proposal about correcting institutional representations of a dispute remains testable. The central Plunkett-Sundell text is still unavailable in full.

The wrong answer to an accurately reported complaint

A resident tells a planning board, “This development is not safe.” The engineer replies, “It satisfies the safety standard.” The resident knows that it satisfies the current numerical standard. Her objection is that the standard excludes an important kind of risk. The engineer understands this and defends the narrower standard as appropriate for licensing development.

Suppose a mediator prepares a cross-language report. It accurately records that the development meets the official threshold and that the resident uses “safe” more stringently. It then presents the disagreement as resolved once these different usages have been explained. No quoted sentence need be mistranslated for the report to misrepresent what remains at issue: whether the official threshold is adequate to authorize the project.

A second report preserves the same descriptive information but adds that the participants disagree about which threshold should govern approval. The difference is not a disagreement about the current measurement. It is a disagreement about the practical authority of a standard.

This is a constructed example. It does not establish that professional translators routinely make this mistake. It isolates a possible defect in a kind of institutional mediation, including translated summaries, intake forms, and automated complaint classifications.

What should a good representation preserve?

The proposed requirement is limited to a representation used as an institution's operative account of a dispute. It should leave the receiving institution able to identify the decision being challenged, the ground of the challenge, and whether resolving a descriptive uncertainty would actually answer it. If the representation instead sends the complaint to a procedure that cannot address its ground, sentence-level accuracy is insufficient for that institutional purpose.

This requirement does not demand preservation of every association or nuance. Nor does it claim that a translation preserving all relevant pragmatic information could still lose the disagreement. That would be contradictory. The point is to specify which pragmatic information counts as relevant when the output will determine what kind of complaint an institution considers.

In the example, an engineer can answer whether the building meets the existing threshold. That answer does not settle whether the threshold should govern permission. A report that turns the second challenge into the first can cause the board to regard an unanswered complaint as answered. The resulting wrong is not simply that the resident dislikes the chosen vocabulary. It is that an institutional act of representation changes what she is treated as having contested.

Where the responsibility falls

It would be unfair to place the entire burden on individual interpreters. They may lack the time, authority, or specialized knowledge to resolve the ambiguity. The institution that adopts the translated summary as authoritative should provide a proportionate way to correct its account of the dispute before using that account to close the matter.

One practical implication is a limited confirmation step: “We understand you to be disputing the measurement, rather than the rule. Is that correct?” The step should identify a genuine ambiguity, not require complainants to approve a document they cannot understand. An interpreter can flag the ambiguity without endorsing either party's substantive position. An authorized decision-maker must then determine where the dispute belongs.

The obligation is especially strong when a foreseeable ambiguity concerns the difference between compliance with a standard and the legitimacy of the standard. It is weaker where the

institution does not rely on the representation to limit any person's opportunities to answer or challenge.

The contribution cannot be the broad diagnosis

Angermeyer's Translation as discrimination (2023) already shows how institutional translation can prioritize institutional needs and position language users as passive recipients rather than participants. Sadler, Baker, and Engebretsen's Translation and the Political (2024) likewise rejects a politically innocent conception of translation. These are substantial predecessors, not ornamental citations.

Plunkett and Sundell's 2013 paper concerns genuine disagreement that need not rest on speakers using an evaluative term with identical meanings. Its abstract is accessible, but the full text was not obtained from the attempted author, repository, or publisher routes. The precise relationship between its account and the present example therefore remains to be checked. This draft does not attribute a complete theory of metalinguistic negotiation to an unread text.

The possible increment is a duty to permit correction of the institution's classification of a dispute, triggered by the role that the mediated representation will play in disposing of the complaint. Even this is likely to overlap with procedural fairness and existing interpreting ethics. It is a candidate for comparison, not an established gap.

Difficult cases

What if the participant changes her complaint after seeing the summary? Correction need not mean unlimited revision without procedural consequences. The institution can distinguish correcting an attribution from introducing a new claim. But it cannot decide that every challenge to its summary is opportunistic simply because the original wording was ambiguous.

What if the dispute is unreasonable? Preserving it does not require accepting it. An institution may reject a challenge to a legitimate standard while accurately identifying that challenge. The right proposed concerns being answered on the relevant ground, not securing the preferred answer.

What if participants disagree about what their disagreement is? A mediator should not manufacture consensus. The operative record may need to preserve competing accounts and identify why the institution adopts one for a particular purpose. The standard is contestable adequacy, not perfect access to speakers' intentions.

The philosophical payoff, if the argument survives

Representation can exercise authority by determining which disagreement officially exists. A focused essay could explain why this creates responsibilities beyond avoiding false translation. To earn originality, however, it must identify a case that existing pragmatic fidelity and procedural fairness requirements do not already settle. Until that comparison succeeds, this remains a promising but unverified development.

References and access

- Angermeyer, Philipp Sebastian. 2023. “Translation as discrimination: Sociolinguistics and inequality in multilingual institutional contexts.” Language in Society. Source. Full HTML available; introduction and institutional interpreting discussion read.
- Sadler, Neil, Mona Baker, and Eivind Engebretsen. 2024. “Translation and the Political: Antagonism and Hegemony.” Translation Studies 17(1):152-168. Source. Author-accepted manuscript introduction and theoretical setup read.
- Plunkett, David, and Tim Sundell. 2013. “Disagreement and the Semantics of Normative and Evaluative Terms.” Philosophers' Imprint 13(23):1-37. Source. Abstract/metadata only; full text requested.

You may still owe me, but I may owe you an explanation

Candidate C045 · Development draft · 22 October 2026

Originality verdict: Best current development prospect among these eight, not a demonstrated original contribution. The candidate is a positive duty of proportionate assistance borne by someone pressing a surviving claim, rather than simply an excuse for the person who has forgotten.

An obligation that survives its memory

Mara borrowed £800 from Ellis and promised repayment. Later, through no fault of her own, she lost the relevant memory. Stipulate that she remains the same person, that the promise remains binding, and that she is presently capable of understanding and discharging it. She has not strategically destroyed records or avoided inquiry. None of these assumptions settles how Ellis should now address her.

Ellis has the original messages and a clear transfer record. Mara remembers neither the loan nor its amount. Ellis tells her that she owes £800. Given the surrounding uncertainty, this bare assertion is insufficient for her reasonably to distinguish the genuine demand from an honest mistake about another transaction. Sending the records would resolve that uncertainty at negligible cost.

It seems insufficient for Ellis merely to refrain from blaming her while continuing to insist on repayment and refusing access to the records. Mara has an interest in recognizing and meeting a real obligation, not merely in escaping criticism. Can that interest ground a claim against Ellis to provide assistance?

The positive duty

The proposed answer is qualified. When a claimant presses performance of an obligation whose basis the other person cannot reasonably recover, the claimant normally owes proportionate help making that basis accessible, if the claimant controls an adequate means of doing so. The duty is strongest where the epistemic obstacle is nonculpable, the requested assistance is inexpensive and safe, and continued insistence otherwise leaves the addressee unable to assess the demand.

“Making accessible” does not mean persuading someone until they agree. Nor does it mean making them relive the original experience. In the loan case it means identifying the transaction and providing available corroboration. Once Ellis has provided adequate evidence, Mara cannot suspend performance merely by reporting that the debt still feels unfamiliar.

The claim also does not make the debt disappear before explanation. Three questions must remain separate: whether a debt exists, whether Mara is blameworthy for failing to pay, and what Ellis owes her while pressing payment. The argument concerns the third question. It does not derive a general rule against every form of enforcement from a rule about blame.

Why refraining from blame is not enough

Suppose Ellis promises never to criticize Mara, privately or publicly, but continues to demand payment, declines to release the records, and leaves her uncertain indefinitely. The absence of blame does not remove the avoidable burden imposed by a live claim she cannot responsibly evaluate. She must either risk failing a real obligation or satisfy a demand without adequate grounds for trusting its basis.

Ellis is specially placed to remove that burden. More importantly, the burden arises in their continuing interaction over Ellis's own claim. This makes the duty more specific than a general requirement to help strangers whenever assistance is cheap. By choosing to press a demand while controlling a proportionate means of assessing it, Ellis acquires responsibility for making the demand practically answerable.

That final step is the argument's vulnerable normative premise. It must be defended, not treated as a consequence of the word “claim.” One defense is that pressing a demand seeks another person's responsible recognition of one's standing. Withholding readily available grounds

needed for that recognition, without a countervailing reason, treats the person's responsible agency as dispensable while still asking them to respond as an obligated agent.

This does not prove that every valid obligation must be consciously intelligible at every moment. It concerns a narrower situation: one person actively seeks compliance from another and can cheaply remove a recognized obstacle to responsible acknowledgment.

Why memory matters, and why it may not be essential

Liao and Sandberg's The Normativity of Memory Modification (2008), §VIII, already examines moral obligations surrounding memory, including the distinction between retaining an experience and preserving relevant information externally. The present proposal cannot be advertised as discovering that memory has moral consequences or that records can substitute for recollection.

Darwall's account of second-personal reasons already connects demands with authority and accountability. A publicly available draft of his opening chapter was consulted; the argument here does not claim to refute or complete his published book. Its possible addition concerns the distribution of epistemic assistance within an existing relation of claim and obligation.

Memory loss is a revealing test case because personal continuity and epistemic continuity come apart. But the same principle may apply when records were lost in a fire or an institutional claim changes hands. If so, the eventual essay should be about responsibilities of claimants under reasonable epistemic dependence. Presenting it solely as a novel ethics of memory would obscure its real scope.

Protecting claimants and victims

The proposal must not require victims to repeatedly justify that they were wronged, expose themselves to danger, or relive traumatic events. The cost and safety conditions are substantive. Testimony can itself be adequate evidence; corroboration is not automatically owed whenever someone demands it. A reasonable institutional intermediary may be appropriate. Sometimes assistance is too burdensome to require at all.

Urgency also matters. Someone causing immediate harm can be required to stop before receiving a historical account of every obligation involved. The loan case concerns a recoverable past commitment and a modest delay, not an emergency.

Strategic forgetting does not dissolve the original obligation or automatically generate an unlimited new entitlement. Even a culpably forgetful person may sometimes need information to comply, but responsibility for the cost can change. A full account should distinguish the duty to supply necessary information from the question who ultimately bears the expense of supplying it.

What would defeat the novelty claim?

The main threat is that ordinary norms of asserting debts, substantiating demands, or facilitating another's performance already contain the whole principle. The crucial comparison is therefore not merely with memory-modification ethics. It is with promissory morality, second-personal authority, and affirmative duties of epistemic assistance. Searches so far have not established that this specific claim is absent there.

The candidate earns further work because it offers a definite target: a positive obligation that remains even when blame is fully suspended, owed by the claimant rather than by an undifferentiated community. If that target survives the next literature comparison, the essay could clarify how recognizing someone's obligation also makes demands on us.

References and access

- Liao, S. Matthew, and Anders Sandberg. 2008. “The Normativity of Memory Modification.” Neuroethics 1:85-99. Source. Author manuscript: Source. Sections VII-VIII and opening of IX read, including the complete moral-obligations section.
- Darwall, Stephen. 2006. The Second-Person Standpoint: Morality, Respect, and Accountability. Harvard University Press. A prepublication draft of the opening chapter, not the entire published book, was consulted at Source. No text from the draft is quoted or redistributed here.

When the evidence needed to reopen a case is behind the closed door

Candidate C085 · Development draft · 22 October 2026

Originality verdict: Merge with C004. The narrower argument concerns an institution that makes access to evidence conditional on already possessing that evidence. Generic defeasibility and the importance of criticism are established.

Closure can be justified

An institution cannot investigate every question forever. Eventually a safety panel must issue a conclusion, an archive must complete its catalogue, or a research group must stop collecting data. Resources are finite, and prolonging one inquiry can prevent more valuable inquiries elsewhere.

The philosophical mistake would be to infer from that practical need that any later challenge may be rejected unless its proponent first meets a demanding evidential threshold. Sometimes the institution controls the very material that would enable the proponent to meet that threshold. A requirement intended to screen weak challenges can then become a mechanism for preventing strong ones from developing.

This is not an objection to Jane Friedman's account of belief as a settling attitude in Inquiry and Belief (2019). In §4.1, she explicitly separates being settled now from whether one will or should inquire later. It would misrepresent her paper to say that it treats belief as permanent immunity from reconsideration. Our question concerns institutions' responsibilities for access and costs, not an alleged defect in that account of individual belief.

A challenge that cannot get started

A safety panel concludes that a bridge component is reliable. Years later a maintenance worker notices an unusual pattern of failures. The panel permits reopening only on evidence of a systematic defect. Yet the relevant production records, necessary to distinguish a systematic defect from ordinary wear, are available only during a formally reopened inquiry.

The worker's observations do not establish that the original conclusion was mistaken. They may eventually prove irrelevant. Nevertheless, requiring proof of the defect as a condition of access creates a circular barrier. The panel cannot fairly defend this barrier merely by saying that unsubstantiated concerns do not merit a complete new investigation.

There is an intermediate possibility: a bounded preliminary check. An independent reviewer could inspect a limited set of records to determine whether the pattern warrants reopening. The worker would have to identify a concrete concern and explain its relevance, but would not have to prove the proposition whose proof requires the withheld material.

The proposed principle

When an institution's control over relevant evidence prevents a reasonably specified challenge from meeting its reopening threshold, the institution normally owes a proportionate route to preliminary assessment that does not presuppose satisfaction of that threshold. Where the institution's own arrangements create the obstacle, it also bears a presumptive responsibility for the additional cost of overcoming it.

This is not a right to full investigation on demand. It is a constraint on the relation between the threshold and the means available for satisfying it. The preliminary stage may legitimately reject a complaint. It must, however, be able to distinguish absence of evidence after an appropriate check from absence produced by refusal to permit the check.

The cost claim requires separate defense. An institution might fairly distribute routine review costs across those participating in a voluntary practice. But it cannot always shift an extraordinary evidential burden onto a person simply because that person was first to articulate a problem created by its own information arrangements. Being the messenger does not make someone the natural insurer of the system's corrigibility.

Why this matters philosophically

The proposal separates the permissibility of stopping work from the authority to impose a continuing conclusion on other people. The panel might have ample justification for its original decision. That historical justification does not settle whether its present procedures treat challengers fairly.

The wrong can therefore exist even if the bridge component is in fact reliable. A challenger need not establish the final truth of the complaint before having a claim to an adequate process for assessing it. Otherwise procedural protection would belong only to those who could already demonstrate that they would win.

This point concerns authority and burden allocation, not a claim that every closed inquiry remains epistemically defective. It allows justified closure while limiting what an institution may demand from those who seek a reasoned reconsideration.

The strongest objections

First, preliminary review itself consumes resources and invites strategic complaints. The response must include thresholds appropriate to that limited stage: identifiable grounds, a connection to the original conclusion, and some prospect that accessible material could bear on the concern. Repeated identical complaints need not repeatedly trigger new review. Where practicable, independent triage can prevent the original decision-makers from being the sole judges of whether a challenge deserves attention.

Second, confidentiality may prevent direct access to the records. Preliminary assessment can sometimes be performed by a trusted reviewer who reports only what is needed for the reopening decision. Where even that would create disproportionate harm, the institution must consider other ways to assess the concern. The principle does not eliminate competing rights.

Third, this sounds like familiar procedural fairness. That may be exactly right. Longino's account already makes effective critical interaction central, and administrative justice has long concerned access to remedies. A full originality claim would need a close comparison with those traditions, particularly their treatment of inaccessible evidence and screening thresholds. This audit has not completed that comparison.

Why it belongs with C004

C004 concerns keeping a future checking route available. C085 concerns allowing an appropriately specified challenger to use it. An archive with no fair access procedure is insufficient; a generous reopening rule with no surviving evidence may also be insufficient. Together the drafts propose responsibilities for making correction possible before and after a challenge arises.

That unity is a reason to merge them. Two headings do not make two discoveries. The combined essay's best chance is to defend a specific responsibility for the burdens an institution itself creates when it makes future correction harder, rather than to announce the familiar ideal of remaining open to criticism.

References and access

- Friedman, Jane. 2019. “Inquiry and Belief.” Noûs 53:296-315. Source. Author accepted manuscript, dated 2017: Source. Introduction, §2 opening, §4.1, and conclusion read. Cite the published 2019 article.
- Longino, Helen E. 2004. “How Values Can Be Good for Science.” In Science, Values, and Objectivity, pp. 127-142. Source. Relevant account of criticism and later reassessment read.
- See C004 for the outstanding comparison with Porat and Stein on evidential damage. That full text remains necessary before claiming the combined principle is new.

A complaint that no department is authorized to hear

Candidate C021 · Development draft · 22 October 2026

Originality verdict: Reject the original broad thesis as a discovery. Standing to contest the normative basis of institutional decisions is already part of the literature. The narrower proposal below concerns responsibility for gaps between complaint procedures; it too remains unverified.

The original claim does not survive

The starting idea was that someone can legitimately contest an institution's objective even when the institution has correctly applied its rule. This is important, but importance does not make it original.

Vredenburgh's The Right to Explanation (2022), particularly §4.1, explicitly connects explanation with input into rules and with their normative reasons. Karlan and Kugelberg's No right to an explanation (2025), §5, likewise discusses the normative reasons supporting institutional decisions. Bayamlıoğlu's Contesting Automated Decisions (2018) advertises scrutiny extending to inference and rulemaking; only its abstract was accessible. The broad candidate should therefore not be submitted as a new discovery about algorithms.

The remaining question is more specific. What happens when an institution nominally permits both individual appeals and collective rule revision, but organizes them so that a relevant challenge fits neither?

Two procedures, one unanswerable complaint

Consider a public housing system that gives priority to applicants expected to find employment quickly. An applicant does not dispute her recorded circumstances or the accuracy with which the rule was applied. She argues that employment potential is an inappropriate objective for allocating this particular housing resource.

The individual appeals office can correct factual and procedural errors but cannot revise the objective. It therefore rejects her complaint as outside its remit. A policy committee can revise objectives but considers only aggregate proposals supported by institution-wide data. The applicant cannot obtain those data because the relevant access procedure requires an active individual appeal.

Each office can truthfully explain its limited competence. Yet their combined arrangement supplies no practical route through which the complaint can receive an answer. Correct application of each office's rules does not establish adequacy of the institution's complaint system.

This is a constructed example, not a report about any existing housing authority. Its interest is the mismatch between the types of reasons a person may have and the types of claims the institution is organized to recognize.

The narrower proposal

An institution that distributes complaint-handling authority across specialized bodies bears responsibility for ensuring that a reasonably specified challenge to its exercise of power has an identifiable route to an appropriately authorized response. A body lacking competence may decline to decide the challenge itself, but its lack of competence does not automatically extinguish the institution's responsibility to address the routing problem.

Sometimes the appropriate response is referral. Sometimes it is a preliminary assessment, access to necessary information, or an explanation of why the matter must be resolved through a different political process. Where no competent route exists, the institution should not represent the failure as if the complainant had simply lost an appeal on the merits.

The claim is not that every office must be empowered to change every rule. Specialization often protects both fairness and competence. The duty falls on the organization of the procedures as a whole, with particular officials assigned responsibility for managing the connections. Otherwise a distributed institution can impose burdens collectively while treating every request for justification as someone else's business.

Why this does not create a personal veto

Standing to seek an answer is different from authority to dictate policy. The applicant's challenge may fail. A legitimate political process may decide that the objective is justified, and an individual officer need not relitigate that decision at every application.

What matters is that the process invoked actually bears on the kind of complaint made. An institution cannot answer a challenge to an objective solely by showing that the objective was accurately implemented. Conversely, it may answer that the objective was adopted through an appropriate process, identify the relevant reasons, and specify the circumstances under which new considerations can be presented.

The proposed duty is compatible with thresholds, time limits, collective representation, and restrictions on repetitive complaints. It is incompatible with treating a mismatch between procedures as evidence that no answerable complaint exists.

Difficulties for the proposal

The first difficulty is identifying the institution. In a distributed public system, no single body may control all the procedures. A duty to repair the gap cannot simply be assigned to whichever official the complainant happens to meet. The essay would need an account of responsibilities attached to design, oversight, and participation, including what limited actors should do when they cannot secure a remedy themselves.

The second difficulty is democratic legitimacy. An administrative body may lack authority to reopen a legislative decision. A referral duty must respect that limit. It may require explaining a political route rather than inventing adjudicative power. Whether the existing political route is adequate is a further substantive question, not something settled by giving a telephone number or a generic link.

The third difficulty is originality. Institutional responsibility, bureaucratic injustice, administrative burdens, and effective remedies are large literatures. A gap between individually proper procedures is not obviously a new philosophical problem. This narrower proposal should be assessed against them before it receives a fresh place on the shortlist.

What to retain

Retain the example and the question of who owes repair when jurisdictional boundaries jointly silence a complaint. Discard the claim that we discovered the difference between contesting inputs and contesting objectives. The philosophical opportunity, if one remains, lies in explaining how an institution acquires responsibilities at the boundaries between its internal divisions that cannot be discharged merely by each division following its own remit.

References and access

- Vredenburgh, Kate. 2022. “The Right to Explanation.” Journal of Political Philosophy 30:209-229. Source. Author manuscript: Source. Especially §4.1 and cost discussion read.
- Karlan, Brett, and Henrik D. Kugelberg. 2025. “No right to an explanation.” Philosophy and Phenomenological Research 111(1):137-156. Source. Introduction and relevant §5 discussion read.
- Bayamlıoğlu, Emre. 2018. “Contesting Automated Decisions: A View of Transparency Implications.” European Data Protection Law Review 4(4):433-446. Source. Institutional abstract read; full text requested.

Source access and requests

Checked 22 October 2026. “Available” means retrievable; “read” identifies the portions actually inspected. Accessible full text is not the same as an exhaustive review.

Six full texts requested from the user

1. **Ariel Porat and Alex Stein (1997), “Liability for Uncertainty: Making Evidential Damage Actionable,” Cardozo Law Review 18:1891-1960.** Record. Needed for C004/C085. Abstract and publication metadata checked; full article unavailable. The earlier repository PDF route failed, and SSRN returned 403 during this pass.

2. **David Plunkett and Tim Sundell (2013), “Disagreement and the Semantics of Normative and Evaluative Terms,” Philosophers' Imprint 13(23):1-37.** Publisher. Needed for C063. Institutional abstract accessible; ANU bitstream timed out, the author-linked copy failed, and Michigan returned 403.

3. **Deborah A. DeMott (2014), “The Fiduciary Character of Agency and the Interpretation of Instructions,” in Philosophical Foundations of Fiduciary Law.** Duke record. Needed for C013. Repository description accessible; linked Duke PDF returned 403.

4. **Philip Pettit (2000), “Democracy, Electoral and Contestatory,” in Designing Democratic Institutions: NOMOS XLII, pp. 105-144.** Author record. Needed for C013/C021. Publication verified; PhilPapers PDF returned 403 and the author-linked document failed.

5. **Emre Bayamlıoğlu (2018), “Contesting Automated Decisions: A View of Transparency Implications,” European Data Protection Law Review 4(4):433-446.** DOI. Needed for C021. Institutional abstract accessible; linked document and SSRN route failed.

6. **André Santos Campos (2026), “Decoupling Generational Sovereignty: The Cross-Temporal Role of the Sovereign People,” Moral Philosophy and Politics, online first, pp. 1-29.** DOI. Needed for C111. Institutional abstract and metadata accessible; institutional PDF and publisher route failed.

If convenient, the **published opening chapter of Stephen Darwall's The Second-Person Standpoint (2006)** would help verify the version consulted for C045. A publicly posted draft was obtained; the final chapter and entire book were not reviewed. This is a version request, not a claim that no relevant Darwall text was accessible.