Elenchus: review edition

Prepared 25 October 2026 for review on 26 October.

The recommended reading set is **Holistic Accuracy and the Lockean Thesis**, **When Acting on Knowledge Makes Free Evidence Harmful**, and **Flattening and Counterfactual Independence**. This is a selection for serious review, not a claim that three original contributions have been established. The first two have the clearest checkable results. The third presents a focused but still disputed interpretation of ordinary counterfactuals.

This edition replaces the older reading packet as the starting point. It incorporates the latest Claude drafts, the earlier review record, and a fresh correction pass. Four manuscripts were substantively revised. The other nine developed tracks and eight exploratory drafts are preserved in the accompanying source archive. Nothing has been submitted to the competition.

Read in this order

1. **Holistic Accuracy and the Lockean Thesis — A5, revision 2.** Start with §3 and then read §§2, 4 and 6. The central example is small enough to check by hand. It challenges the inference from order-independent, truth-directed accuracy maximization to a fixed belief threshold. This is the strongest candidate in this review edition.
2. **When Acting on Knowledge Makes Free Evidence Harmful — A6, revision 2.** Start with the second-weighing example in §3 and the two boundary statements in §4. The proposed contribution is a precise condition for safety across all decision problems, not the familiar observation that treating defeasible knowledge as certainty can cause trouble.
3. **Flattening and Counterfactual Independence — E3, revision 4.** Read the Reunion/Avoidance pair in §3 before the formal material. The revised §5.1 is essential: the linguistic case remains conditional on a stable scenario-relative reading. This is the most vulnerable of the three to a substantive objection rather than an algebraic error.

The combined reading PDF contains this guide followed by those three complete essays. Separate PDFs are supplied for all three and for the causation reserve. The source archive includes editable Markdown, the assessment, methodology records, revision differences, and calculation outputs.

Assessment of the three

A5: a specific counterexample, not the first non-Lockean accuracy theory

What survives. A fixed score over all belief states adds a premium for an error-free record. At credences (0.70, 0.15, 0.15), its unique optimum believes the first world; at (0.75, 0.01, 0.24), its unique optimum does not. The score, option space and proposition are unchanged. The expected accuracies and runner-up margins were reproduced exactly.

The general positive-premium result has a proof, and its sampled parameter cases passed again.

Why it could matter. If the target's minimal premises permit this score, accuracy maximization does not by itself settle the Lockean form of the relation between credence and belief. The dispute must concern further restrictions on epistemic value or rational choice.

Closest predecessors. Dorst's published §4.2 and proof of Proposition 5 were checked. Wang's full 2025 article was recovered; it already treats general, potentially nonadditive utility on deductively cogent belief states. Rothschild's scoring approach and Holguín's nonmonotonicity argument are also acknowledged. The proposed increment is the unrestricted-option counterexample to this particular derivation and the characterization of the record score. Broad claims to have discovered non-Lockean accuracy maximization were removed.

Strongest objection. Sequence-Irrelevance might be intended to require more than order-independent assessment of final states. A defender could demand optimal local choice at every stage, or impose a restriction on which holistic scores count as accuracy. The revision states that interpretive issue and its explicit permission bridge. A small example with correct arithmetic does not settle the interpretation.

Changes. Removed the assertion that every valid route to Lockeanism requires separability. ME plus BI is now presented as one sufficient repair. Corrected the unqualified claim that record-score belief is always consistent. Removed the separate split-invariance theorem and its incompletely established minimality claim. Added the representation-dependence objection. Replaced the abstract-only Wang comparison and old Dorst pagination.

Verdict. Advance for external philosophical review. The restricted claim is promising; the assertion that the target's stated premises are all met requires especially close scrutiny. No priority certification.

A6: an exact boundary with important domain restrictions

What survives. In finite statist, comparable K-normality models, a given question is safe for every decision problem exactly when it has no relative defeat. With the specified questions available, every question is safe exactly when the current ken is gapped. These are different quantifiers. The construction of a three-act counterexample to the stated conditional expected-utility repair also survives the finite checks.

Why it could matter. This gives a precise constraint on a proposed way of making knowledge guide numerical choice. It separates cases where the decision rule can guarantee the value of free checking from cases where some decision problem defeats that guarantee.

Closest predecessors. Goodman and Salow already supply the normality framework and defeat examples. Salow already connects dogmatism with Preservation. Hong's full 2026 paper was recovered and already challenges action guidance using normality-based knowledge. The revised contribution claim is the finite safety boundary and the specified

repair's counterexample, not a first discovery that normality can make action guidance problematic.

Strongest objection. The universal boundary ranges over a generous class of questions and stakes. It does not establish irrational avoidance for every ordinary decision problem or refute every knowledge-first decision theory. The continuous measurement example needs a probability model and a positive-measure argument; a finite-state singleton proof cannot do that work.

Changes. Corrected the empty-question value from "always zero" to "nonnegative," with an explicit two-state counterexample to equality. Corrected the existential quantifier over unsafe questions. Restricted the ten-to-one claim to normal situations. Added a positive-measure argument for Gaussian measurement and identified its conditional prior assumption. Removed an unsupported assertion about generic absence of payoff ties and narrowed the encroachment conclusion. Updated the Hong comparison.

Verdict. Advance alongside A5. Its formal boundary is more secure than its broad philosophical rhetoric was. The revised domain and question-class restrictions are part of the result, not qualifications to omit from a later abstract.

E3: a discriminating case whose interpretation remains unsettled

What survives. The explicit Reunion and Avoidance models agree on all first-degree conditionals at actuality while giving different verdicts to the relevant nested conditional. The controls and finite conservativity construction were rerun successfully. Their correctness does not establish that ordinary speakers interpret the sentences that way.

Why it could matter. A stable scenario-relative reading would challenge the natural-language motivation for Flattening while leaving Dorr and Mandelkern's mathematical characterization intact. The example asks what an embedded supposition inherits from the scenario reached by the outer supposition.

Closest predecessor. Dorr and Mandelkern already know that similarity semantics permits formal countermodels and discuss one themselves. Their discussion in the repository's source extract was checked. The potentially new contribution is the independence/avoidance contrast and its diagnostic controls, not the mere existence of a model where Flattening fails. The external O3 manuscript is another treatment of the same project, not independent novelty evidence.

Strongest objection. A temporary context shift could produce the apparent contrast while preserving Flattening for a fixed operator. Different background stories can shift interpretation without changing sentence wording. The draft's earlier attempt to rule this out was unsuccessful. Locality in Buttons likewise needs a counterfactual selection premise; relativity alone does not fix it.

Changes. Rewrote the abstract, thesis, context-shift reply and conclusion. Removed the assertion that no natural-language interpretation could save Flattening. Stated the local-change premise explicitly and restricted the conservativity observation to the finite models actually used. The paper now offers a testable challenge, not a completed empirical refutation.

Verdict. Provisional third reading candidate. Do not treat its third-place inclusion as evidence that its linguistic claim has been independently confirmed. E1 or E2 remains an alternative if this dispute proves fatal to the proposed contribution.

Causation reserve: A4 revision 2

The frozen-witness lemma previously assumed that an endogenous member of a conjunctive cause had unchanged parents. Another member of the conjunction can be its ancestor, so that inference fails. The revised proof uses acyclicity and AC3 minimality: an equation-conforming witness is already generated by intervening on the root members, making any extra endogenous candidate redundant.

The other main correction concerns scope. Four specified uniform default codings are not every possible uniform coding, and a local componentwise filter is not the whole Halpern–Hitchcock framework. The text now distinguishes those restrictions. The locust/switch contrast is defended conditionally rather than described as indisputable.

The existing case program ran again. The repaired lemma is also supported by its written argument; the finite program is not a proof for all causal models. The manuscript remains a reserve because the philosophical objection can be answered by enriching the representation, rejecting a verdict, or leaving the restricted ranking class. Halpern's full 2016 book remains an incomplete comparison. Its importance does not automatically establish the novelty of this restricted impossibility result.

Inventory and dispositions

The project contains **13 developed tracks, approximately 10 distinct arguments, and eight exploratory drafts**. Revisions, model cards, checks and reading packets are not additional essays. "OA" below distinguishes external OpenAI manuscripts from Claude manuscripts with the same A-code.

Track	Current review source	Disposition
Claude A5	Holistic Accuracy, v2 in this edition	First reading candidate
Claude A6	Free Evidence, v2 in this edition	Second reading candidate
E3	Flattening, v4 in this edition	Provisional third candidate
Claude A4	Normality and Actual Causation, v2 here	Corrected reserve
E1	Gettier Knowledge across Independent Sources, v3	Reserve; Gettier diagnosis and increment beyond the target's own model remain decisive

Track	Current review source	Disposition
E2	Ignorable Learning, external v5	Reserve; newest version addresses Gallow and Reflection scope, but needs a fresh full referee decision
O1	Independent Sources and Correlations	Alternate E1 treatment, not a separate discovery
O2	Bayes and the Benchmark	Alternate E2 treatment, not a separate discovery
O3	Independent Interventions	Alternate E3 treatment, not a separate discovery
O4	Evidence and Description, v2	Do not advance as a refutation of a claim its target does not make
O5	How Regret Selects a Learning Standard, v3	Reserve; classification is narrower than generic minimax/score-dependence claims
OA-A4	Finite relay criterion, v2	Research note; no established full apt counterexample
OA-A6	Almost-sure recovery and cost, v0.2	Research note; adequate epistemic realization still missing

The last nine tracks are archived as received, not silently certified by this pass. Claude's earlier negative reviews remain relevant; later revisions are not automatically cleared by answering some objections. The previous E1/E2/E3 preparation set was provisional as well.

The eight exploratory ideas remain: usable revocation; models and accountability; preserving correction; inherited authority; translating disagreement; forgotten obligations; reopening inquiry; and contesting objectives. Their existing audit demotes broad novelty claims and recommends merging preserving correction with reopening inquiry. C045, on claimant assistance after memory loss, remains a research prospect rather than a finished contest entry. The source archive preserves that audit and its outstanding text requests. Further expansion of this pool was not necessary to finish a useful review edition.

Verification and its limits

Fresh checks reproduced A5's exact optima, nine robustness cases, and the core characterization on 887 sampled instances. A6's program examined 64,004 comparable situation–question pairs through four states: all 3,306 relative-defeat cases yielded the constructed negative-value examples, and no sampled KDT failure occurred without relative defeat. Gaussian integration reproduced approximately -0.0499, -0.1202 and

−0.2739 for the displayed 10:1, 20:1 and 100:1 bets. The empty-question correction gives +0.5 in its explicit example. E3's case, control, conservativity and crashing-fragment checks passed. A4's case program completed and is retained.

These are targeted reruns of the existing calculations, not a new software project. The optional Z3 dependency was unavailable in this runtime. E3's historical solver claim was not rerun; its mathematical argument and dependency-free checks were inspected instead. Sampling does not prove general claims, and computation does not prove originality or philosophical significance. Detailed outputs and limitations are in VALIDATION.md.

Source access and review standard

Dorst's published article, Wang's full article, and Hong's full article were accessible in this pass. Those earlier access gaps are closed for the comparisons reported here. Dorr and Mandelkern's relevant discussion was checked through the repository's text extract. No source was treated as read merely because a search result named it.

One additional comparison located in this pass remains inaccessible: **Ginger Schultheis, "Reasonable Inferences for Counterfactuals"** (manuscript, PhilArchive SCHRIF-3). The indexed description concerns reasonable counterfactual inference; its appearance under a Flattening search is not evidence that it contains this result. A PDF would permit the comparison. It is recorded as an unresolved literature lead, not an alleged refutation. The older reserve-track requests, including the published Darwall material and Halpern's book, remain in SOURCE-REQUESTS.md. None prevents reading the completed packet tomorrow.

The assessment applies four questions: what exactly follows; which premise could reasonably be rejected; what the nearest predecessor already establishes; and what philosophical belief would change if the result were accepted. These are synthetic adversarial tests by the revising assistant. No philosopher, contest judge, future model or independent external reviewer has endorsed this edition.

This is the stopping point for this review edition. The unresolved questions are documented rather than hidden behind extra word count or a simulated score. The next substantive judgment belongs to an independent review of the narrowed claims.